

THE DEFINITIVE MASTER REFERENCE HANDBOOK FOR MACHINE LEARNING, DEEP LEARNING & COMPUTER VISION

A Comprehensive One-Stop Pedagogical Guide to 66 Scikit-Learn Algorithms, Mathematical Formulations, Loss Functions, Error Metrics, Computer Vision Architectures, and Unsupervised Learning

Edition 1.0 | Complete Learning & Professional Engineering Manual
Includes Full Mathematical Derivations, Mental Models & Architecture Flowcharts

Contents

List of Abbreviations & Acronyms Index	7
Chapter 1: Foundational Framework of Model Evaluation & Error Metrics	10
1.1 Confusion Matrix Layout & Formulations	10
1.2 Situation-to-Metric Selection Guide	10
1.3 Classification Loss Functions & Metrics Summary	10
1.4 Regression Loss Functions & Evaluation Metrics Summary	11
Chapter 2: Classification Models & Architectures (30 Models)	13
1. LinearSVC	13
Mental Model	13
Mathematical Formulation & Prediction	13
Optimization & Loss Function	13
Meaning Breakdown & Mental Flowchart	13
2. SGDClassifier	13
Mental Model	13
Mathematical Formulation & Prediction	13
Optimization & Loss Function	13
Meaning Breakdown & Mental Flowchart	13
3. MLPClassifier	14
Mental Model	14
Mathematical Formulation & Prediction	14
Optimization & Loss Function	14
Meaning Breakdown & Mental Flowchart	14
4. Perceptron	14
Mental Model	14
Mathematical Formulation & Prediction	14
Optimization & Loss Function	14
Meaning Breakdown & Mental Flowchart	15
5. LogisticRegression	15
Mental Model	15
Mathematical Formulation & Prediction	15
Optimization & Loss Function	15
Meaning Breakdown & Mental Flowchart	15
6. LogisticRegressionCV	15
Mental Model	15
Mathematical Formulation & Prediction	15
Optimization & Loss Function	16
Meaning Breakdown & Mental Flowchart	16
7. SVC (Support Vector Classifier)	16
Mental Model	16
Mathematical Formulation & Prediction	16
Optimization & Loss Function	16
Meaning Breakdown & Mental Flowchart	16
8. CalibratedClassifierCV	16
Mental Model	17
Mathematical Formulation & Prediction	17
Optimization & Loss Function	17
Meaning Breakdown & Mental Flowchart	17
9. PassiveAggressiveClassifier	17
Mental Model	17
Mathematical Formulation & Prediction	17
Optimization & Loss Function	17

Meaning Breakdown & Mental Flowchart	17
10. LabelPropagation	17
Mental Model	18
Mathematical Formulation & Prediction	18
Optimization & Loss Function	18
Meaning Breakdown & Mental Flowchart	18
11. LabelSpreading	18
Mental Model	18
Mathematical Formulation & Prediction	18
Optimization & Loss Function	18
Meaning Breakdown & Mental Flowchart	18
12. RandomForestClassifier	18
Mental Model	19
Mathematical Formulation & Prediction	19
Optimization & Loss Function	19
Meaning Breakdown & Mental Flowchart	19
13. GradientBoostingClassifier	19
Mental Model	19
Mathematical Formulation & Prediction	19
Optimization & Loss Function	19
Meaning Breakdown & Mental Flowchart	19
14. QuadraticDiscriminantAnalysis	20
Mental Model	20
Mathematical Formulation & Prediction	20
Optimization & Loss Function	20
Meaning Breakdown & Mental Flowchart	20
15. HistGradientBoostingClassifier	20
Mental Model	20
Mathematical Formulation & Prediction	20
Optimization & Loss Function	20
Meaning Breakdown & Mental Flowchart	20
16. RidgeClassifierCV	21
Mental Model	21
Mathematical Formulation & Prediction	21
Optimization & Loss Function	21
Meaning Breakdown & Mental Flowchart	21
17. RidgeClassifier	21
Mental Model	21
Mathematical Formulation & Prediction	21
Optimization & Loss Function	21
Meaning Breakdown & Mental Flowchart	21
18. AdaBoostClassifier	22
Mental Model	22
Mathematical Formulation & Prediction	22
Optimization & Loss Function	22
Meaning Breakdown & Mental Flowchart	22
19. ExtraTreesClassifier	22
Mental Model	22
Mathematical Formulation & Prediction	22
Optimization & Loss Function	22
Meaning Breakdown & Mental Flowchart	23
20. KNeighborsClassifier	23
Mental Model	23
Mathematical Formulation & Prediction	23

Optimization & Loss Function	23
Meaning Breakdown & Mental Flowchart	23
21. BaggingClassifier	23
Mental Model	23
Mathematical Formulation & Prediction	23
Optimization & Loss Function	23
Meaning Breakdown & Mental Flowchart	24
22. BernoulliNB	24
Mental Model	24
Mathematical Formulation & Prediction	24
Optimization & Loss Function	24
Meaning Breakdown & Mental Flowchart	24
23. LinearDiscriminantAnalysis	24
Mental Model	24
Mathematical Formulation & Prediction	24
Optimization & Loss Function	24
Meaning Breakdown & Mental Flowchart	24
24. GaussianNB	25
Mental Model	25
Mathematical Formulation & Prediction	25
Optimization & Loss Function	25
Meaning Breakdown & Mental Flowchart	25
25. NuSVC	25
Mental Model	25
Mathematical Formulation & Prediction	25
Optimization & Loss Function	25
Meaning Breakdown & Mental Flowchart	25
26. DecisionTreeClassifier	26
Mental Model	26
Mathematical Formulation & Prediction	26
Optimization & Loss Function	26
Meaning Breakdown & Mental Flowchart	26
27. NearestCentroid	26
Mental Model	26
Mathematical Formulation & Prediction	26
Optimization & Loss Function	26
Meaning Breakdown & Mental Flowchart	26
28. ExtraTreeClassifier	27
Mental Model	27
Mathematical Formulation & Prediction	27
Optimization & Loss Function	27
Meaning Breakdown & Mental Flowchart	27
29. CheckingClassifier	27
Mental Model	27
Mathematical Formulation & Prediction	27
Optimization & Loss Function	27
Meaning Breakdown & Mental Flowchart	27
30. DummyClassifier	27
Mental Model	27
Mathematical Formulation & Prediction	27
Optimization & Loss Function	28
Meaning Breakdown & Mental Flowchart	28

Chapter 3: Regression Models & Architectures (36 Models) 29

31. SVR	29
Mental Model	29
Mathematical Formulation & Prediction	29
Optimization & Loss Function	29
Meaning Breakdown & Mental Flowchart	29
32. RandomForestRegressor	29
Mental Model	29
Mathematical Formulation & Prediction	29
Optimization & Loss Function	29
Meaning Breakdown & Mental Flowchart	29
33. ExtraTreesRegressor	30
Mental Model	30
Mathematical Formulation & Prediction	30
Optimization & Loss Function	30
Meaning Breakdown & Mental Flowchart	30
34. AdaBoostRegressor	30
Mental Model	30
Mathematical Formulation & Prediction	30
Optimization & Loss Function	30
Meaning Breakdown & Mental Flowchart	30
35. NuSVR	30
Mental Model	30
Mathematical Formulation & Prediction	31
Optimization & Loss Function	31
Meaning Breakdown & Mental Flowchart	31
36. GradientBoostingRegressor	31
Mental Model	31
Mathematical Formulation & Prediction	31
Optimization & Loss Function	31
Meaning Breakdown & Mental Flowchart	31
37. KNeighborsRegressor	31
Mental Model	31
Mathematical Formulation & Prediction	31
Optimization & Loss Function	32
Meaning Breakdown & Mental Flowchart	32
38. HistGradientBoostingRegressor	32
Mental Model	32
Mathematical Formulation & Prediction	32
Optimization & Loss Function	32
Meaning Breakdown & Mental Flowchart	32
39. BaggingRegressor	32
Mental Model	32
Mathematical Formulation & Prediction	32
Optimization & Loss Function	32
Meaning Breakdown & Mental Flowchart	33
40. MLPRegressor	33
Mental Model	33
Mathematical Formulation & Prediction	33
Optimization & Loss Function	33
Meaning Breakdown & Mental Flowchart	33
41. HuberRegressor	33
Mental Model	33
Mathematical Formulation & Prediction	33
Optimization & Loss Function	33

	Meaning Breakdown & Mental Flowchart	34
42.	LinearSVR	34
	Mental Model	34
	Mathematical Formulation & Prediction	34
	Optimization & Loss Function	34
	Meaning Breakdown & Mental Flowchart	34
43.	RidgeCV	34
	Mental Model	34
	Mathematical Formulation & Prediction	34
	Optimization & Loss Function	34
	Meaning Breakdown & Mental Flowchart	34
44.	BayesianRidge	35
	Mental Model	35
	Mathematical Formulation & Prediction	35
	Optimization & Loss Function	35
	Meaning Breakdown & Mental Flowchart	35
45.	Ridge	35
	Mental Model	35
	Mathematical Formulation & Prediction	35
	Optimization & Loss Function	35
	Meaning Breakdown & Mental Flowchart	35
46.	LinearRegression	36
	Mental Model	36
	Mathematical Formulation & Prediction	36
	Optimization & Loss Function	36
	Meaning Breakdown & Mental Flowchart	36
47.	TransformedTargetRegressor	36
	Mental Model	36
	Mathematical Formulation & Prediction	36
	Optimization & Loss Function	36
	Meaning Breakdown & Mental Flowchart	36
48.	LassoCV	37
	Mental Model	37
	Mathematical Formulation & Prediction	37
	Optimization & Loss Function	37
	Meaning Breakdown & Mental Flowchart	37
49.	ElasticNetCV	37
	Mental Model	37
	Mathematical Formulation & Prediction	37
	Optimization & Loss Function	37
	Meaning Breakdown & Mental Flowchart	37
50.	LassoLarsCV	37
	Mental Model	38
	Mathematical Formulation & Prediction	38
	Optimization & Loss Function	38
	Meaning Breakdown & Mental Flowchart	38
51.	LassoLarsIC	38
	Mental Model	38
	Mathematical Formulation & Prediction	38
	Optimization & Loss Function	38
	Meaning Breakdown & Mental Flowchart	38
52.	LarsCV	38
	Mental Model	38
	Mathematical Formulation & Prediction	39

	Optimization & Loss Function	39
	Meaning Breakdown & Mental Flowchart	39
53. Lars		39
	Mental Model	39
	Mathematical Formulation & Prediction	39
	Optimization & Loss Function	39
	Meaning Breakdown & Mental Flowchart	39
54. SGDRegressor		39
	Mental Model	39
	Mathematical Formulation & Prediction	39
	Optimization & Loss Function	39
	Meaning Breakdown & Mental Flowchart	40
55. RANSACRegressor		40
	Mental Model	40
	Mathematical Formulation & Prediction	40
	Optimization & Loss Function	40
	Meaning Breakdown & Mental Flowchart	40
56. ElasticNet		40
	Mental Model	40
	Mathematical Formulation & Prediction	40
	Optimization & Loss Function	40
	Meaning Breakdown & Mental Flowchart	41
57. Lasso		41
	Mental Model	41
	Mathematical Formulation & Prediction	41
	Optimization & Loss Function	41
	Meaning Breakdown & Mental Flowchart	41
58. OrthogonalMatchingPursuitCV		41
	Mental Model	41
	Mathematical Formulation & Prediction	41
	Optimization & Loss Function	41
	Meaning Breakdown & Mental Flowchart	41
59. ExtraTreeRegressor		42
	Mental Model	42
	Mathematical Formulation & Prediction	42
	Optimization & Loss Function	42
	Meaning Breakdown & Mental Flowchart	42
60. PassiveAggressiveRegressor		42
	Mental Model	42
	Mathematical Formulation & Prediction	42
	Optimization & Loss Function	42
	Meaning Breakdown & Mental Flowchart	42
61. GaussianProcessRegressor		43
	Mental Model	43
	Mathematical Formulation & Prediction	43
	Optimization & Loss Function	43
	Meaning Breakdown & Mental Flowchart	43
62. OrthogonalMatchingPursuit		43
	Mental Model	43
	Mathematical Formulation & Prediction	43
	Optimization & Loss Function	43
	Meaning Breakdown & Mental Flowchart	43
63. DecisionTreeRegressor		44
	Mental Model	44

Mathematical Formulation & Prediction	44
Optimization & Loss Function	44
Meaning Breakdown & Mental Flowchart	44
64. DummyRegressor	44
Mental Model	44
Mathematical Formulation & Prediction	44
Optimization & Loss Function	44
Meaning Breakdown & Mental Flowchart	44
65. LassoLars	45
Mental Model	45
Mathematical Formulation & Prediction	45
Optimization & Loss Function	45
Meaning Breakdown & Mental Flowchart	45
66. KernelRidge	45
Mental Model	45
Mathematical Formulation & Prediction	45
Optimization & Loss Function	45
Meaning Breakdown & Mental Flowchart	45
Chapter 4: Computer Vision & Deep Learning Mental Models	46
4.1 Classic CNNs & Vision Transformers	46
4.2 Object Detection Architectures	46
4.3 Image Segmentation Architectures	47
Chapter 5: Unsupervised Learning & Vector Mathematics	49
5.1 Distance & Similarity Metrics	49
5.2 Clustering Algorithms	49
5.3 Dimensionality Reduction Techniques	50

List of Abbreviations & Acronyms Index

Abbreviation	Full Term	Brief Description
AIC	Akaike Information Criterion	Model selection score balancing likelihood and parameter count
AUC	Area Under the Curve	Aggregate measure of performance across all classification thresholds
BCE	Binary Cross-Entropy	Log loss metric for binary 0/1 probability classification
BIC	Bayesian Information Criterion	Model selection criterion with stronger penalty for parameters than AIC
CCE	Categorical Cross-Entropy	Multi-class log loss over Softmax probability distributions
CNN	Convolutional Neural Network	Deep network using spatial convolutions for visual grid processing
CV	Cross-Validation	Resampling procedure to evaluate ML model generalization
DBSCAN	Density-Based Spatial Clustering	Clustering algorithm based on dense neighborhood connectivity
DETR	DEtection TRansformer	End-to-end object detector using Transformer encoder-decoder
FCN	Fully Convolutional Network	CNN architecture for pixel-by-pixel semantic segmentation

Abbreviation	Full Term	Brief Description
GMM	Gaussian Mixture Model	Probabilistic soft clustering model assuming Gaussian components
GP	Gaussian Process	Non-parametric Bayesian probabilistic regression framework
IoU	Intersection over Union	Overlap evaluation metric for bounding boxes and masks
KDTree	k-Dimensional Tree	Space-partitioning data structure for fast nearest neighbor search
KL	Kullback-Leibler	Relative entropy measuring divergence between probability distributions
KNN	K-Nearest Neighbors	Instance-based non-parametric classifier/regressor
LARS	Least Angle Regression	Stepwise forward feature selection algorithm
LDA	Linear Discriminant Analysis	Linear supervised classification & dimensionality reduction
LOOCV	Leave-One-Out Cross-Validation	Cross-validation where fold size equals 1 sample
MAE	Mean Absolute Error	Average absolute error distance $ y - \hat{y} $
MAPE	Mean Absolute Percentage Error	Scale-independent relative error percentage
MCC	Matthews Correlation Coefficient	Balanced correlation coefficient for binary classification
MLP	Multi-Layer Perceptron	Classical feed-forward artificial neural network
MSE	Mean Squared Error	Average squared prediction error distance $(y - \hat{y})^2$
MSLE	Mean Squared Logarithmic Error	Relative logarithmic squared error metric
OLS	Ordinary Least Squares	Linear regression minimizing sum of squared residuals
OMP	Orthogonal Matching Pursuit	Greedy sparse signal approximation algorithm
PCA	Principal Component Analysis	Unsupervised orthogonal variance dimensionality reduction
QDA	Quadratic Discriminant Analysis	Generative classifier with class-specific covariance matrices
RBF	Radial Basis Function	Gaussian kernel measuring Euclidean distance similarity
RMSE	Root Mean Squared Error	Square root of MSE; error in original target units
ROC	Receiver Operating Characteristic	Curve plotting True Positive Rate vs False Positive Rate
RANSAC	RANdom SAMple Consensus	Iterative robust regressor resistant to heavy outliers
SGD	Stochastic Gradient Descent	Iterative optimization using mini-batch sample gradients
sMAPE	Symmetric MAPE	Symmetric relative percentage error metric bounded [0%, 200%]
SMO	Sequential Minimal Optimization	Algorithm for fast quadratic optimization in SVMs
SSD	Single Shot MultiBox Detector	Real-time one-stage object detection CNN

Abbreviation	Full Term	Brief Description
SVC / SVR	Support Vector Classifier / Regressor	Maximum-margin hyperplane support vector machine
t-SNE	t-Distributed Stochastic Neighbor Embedding	Non-linear manifold visualization technique
UMAP	Uniform Manifold Approximation & Projection	Fast non-linear manifold dimensionality reduction
ViT	Vision Transformer	Transformer architecture applied directly to image patches
WAPE	Weighted Absolute Percentage Error	Volume-weighted absolute percentage error metric
YOLO	You Only Look Once	Real-time single-stage object detector

Chapter 1: Foundational Framework of Model Evaluation & Error Metrics

1.1 Confusion Matrix Layout & Formulations

	Actual Positive	Actual Negative
Predicted Positive	TP (True Positive) Correctly identified positive class	FP (False Positive) Type I Error - False Alarm
Predicted Negative	FN (False Negative) Type II Error - Missed Case	TN (True Negative) Correctly identified negative class

1.2 Situation-to-Metric Selection Guide

Business / Modeling Situation	Recommended Metric	Key Reason / Focus
Classes are balanced	Accuracy	Overall ratio of correct predictions
False positives are expensive	Precision	Minimizes false alarms (e.g., spam filter)
False negatives are expensive	Recall / Sensitivity	Minimizes missed critical cases (e.g., medical diagnosis)
Need balance of Precision + Recall	F1 Score	Harmonic mean balancing FP and FN penalties
Very imbalanced dataset	F1 / MCC / PR-AUC	Avoids accuracy inflation on majority class
Need probability quality	Log Loss	Penalizes confident incorrect probability predictions
Need threshold-independent ranking	ROC-AUC	Measures class separation ability across all thresholds
Medical diagnosis	Recall + Specificity + ROC-AUC	Ensures high detection while tracking true negatives
Image/object classification	Accuracy + F1 + Confusion Matrix	Comprehensively tracks per-class errors
Multi-class classification	Macro F1 + Weighted F1 + Accuracy	Balances overall performance and minority classes

1.3 Classification Loss Functions & Metrics Summary

Function Name	Full Name / Variant	Main Purpose	Key Property & Formula Note
BCE	Binary Cross Entropy	Binary classification	Penalizes wrong probability outputs for 0/1 target
BCE with Logits	Binary CE with Logits	Binary classification directly from raw logits	More numerically stable (combines Sigmoid + BCE)
Categorical CE	Categorical Cross Entropy	Multi-class classification	Requires one-hot encoded targets with Softmax
Sparse Categorical CE	Sparse Categorical CE	Multi-class classification	Accepts integer class labels directly (memory efficient)
Focal Loss	Focal Loss	Handles severe class imbalance	Down-weights easy examples with factor $(1-p)^{\gamma}$
Hinge Loss	Hinge Loss	SVM-style classification	Maximizes margin $\max(0, 1 - y \cdot f(x))$ for $y \in \{-1, +1\}$
Squared Hinge Loss	Squared Hinge Loss	SVM-style classification	Quadratic penalty for margin violations $[\max(0, 1 - y \cdot f(x))]^2$

Function Name	Full Name / Variant	Main Purpose	Key Property & Formula Note
KL Divergence	Kullback-Leibler Divergence	Probability distribution fitting	Measures divergence between predicted Q and true P
Label Smoothing Loss	Label Smoothing CE	Prevents overconfident predictions	Softens hard 0/1 targets into 0.05/0.95 distributions

1.4 Regression Loss Functions & Evaluation Metrics Summary

Metric / Loss	Full Name	Range	Best	Main Purpose & Key Property
MAE	Mean Absolute Error	$0-\infty$	$\downarrow 0$	Average absolute error. Robust to outliers.
MedAE	Median Absolute Error	$0-\infty$	$\downarrow 0$	Median prediction error. Highly robust against extreme outliers.
MSE	Mean Squared Error	$0-\infty$	$\downarrow 0$	Penalizes large errors strongly (squared error).
RMSE	Root Mean Squared Error	$0-\infty$	$\downarrow 0$	Square root of MSE; error is in same units as target.
R^2	R-squared (Coeff. of Det.)	$-\infty$ to 1	$\uparrow 1$	Proportion of target variance explained by model.
Adjusted R^2	Adjusted R-squared	$-\infty$ to 1	$\uparrow 1$	R^2 penalized for adding non-informative feature variables.
Explained Var.	Explained Variance Score	$-\infty$ to 1	$\uparrow 1$	Measures variation captured (ignores mean bias).
MAPE	Mean Absolute % Error	$0-\infty\%$	$\downarrow 0\%$	Percentage-based error. Problematic when actual $y \approx 0$.
sMAPE	Symmetric MAPE	$0-200\%$	$\downarrow 0\%$	Symmetric percentage error; less sensitive to zero target issues.
MSPE	Mean Squared % Error	$0-\infty$	$\downarrow 0$	Squared percentage error for severe relative penalty.
MSLE	Mean Squared Log Error	$0-\infty$	$\downarrow 0$	Penalizes relative log-scale error; useful for skewed targets.
RMSLE	Root Mean Squared Log Error	$0-\infty$	$\downarrow 0$	Root version of MSLE; reduces impact of large target scale.

Metric / Loss	Full Name	Range	Best	Main Purpose & Key Property
Huber Loss	Huber / Smooth L1 Loss	$0-\infty$	$\downarrow 0$	Combines MSE (small errors) + MAE (large errors). Outlier-robust.
Log-Cosh Loss	Logarithm of Hyperbolic Cosine	$0-\infty$	$\downarrow 0$	Smooth robust loss, twice-differentiable everywhere.
Quantile Loss	Pinball Loss	$0-\infty$	$\downarrow 0$	Predicts specified percentiles (used for prediction intervals).
Max Error	Maximum Error	$0-\infty$	$\downarrow 0$	Captures the worst single prediction error in dataset.

Chapter 2: Classification Models & Architectures (30 Models)

1. LinearSVC

Mental Model

Find a linear hyperplane that separates classes with the largest possible geometric margin.

Mathematical Formulation & Prediction

$$f(x) = \hat{w}^T \cdot x + b$$

Prediction: $y_hat = +1$ if $f(x) > 0$ else -1

Optimization & Loss Function

$$\min_{\{w,b\}} 0.5 * ||w||^2 + C * \sum \max(0, 1 - y_i * (\hat{w}^T * x_i + b))^2$$

The loss term is Squared Hinge Loss (or standard Hinge Loss).

Meaning Breakdown & Mental Flowchart

$$||w||^2$$

↓

Keep boundary simple / maximize geometric margin

hinge loss

↓

Punish incorrectly classified or insufficiently separated samples

C hyperparameter

↓

Trade-off between wide margin and classification error

2. SGDClassifier

Mental Model

Take a small sample or single data point, calculate the prediction error, and instantly update weights via gradient step.

Mathematical Formulation & Prediction

$$\text{General Objective: } J(w) = (1/n) * \sum L(y_i, f(x_i)) + \lambda * R(w)$$

$$\text{Single Step Update: } w_{\{t+1\}} = w_t - \eta * \nabla J(w_t)$$

Optimization & Loss Function

Configurable loss $L(y, f(x))$:

- L = Hinge (SVM)
- L = Log loss (Logistic Regression)
- L = Modified Huber (probabilistic SVM)

Regularization $R(w)$: L2 norm $0.5 * ||w||^2$, L1 norm $||w||_1$, or ElasticNet.

Meaning Breakdown & Mental Flowchart

prediction

↓

Make quick output guess for current sample

calculate error

↓

Compute loss value $L(y_i, y_hat_i)$

calculate gradient

↓

Compute gradient ∇J with respect to weight vector w

change weights

↓

Update weights $w_{t+1} = w_t - \eta * \nabla J$ and repeat

3. MLPClassifier

Mental Model

Stack multiple layers of artificial neurons to learn complex, non-linear feature representations automatically.

Mathematical Formulation & Prediction

Single Neuron: $z = w^T * x + b$, **Activation:** $a = f(z)$

Network Layers: $h_1 = f(W_1 * x + b_1)$, $h_2 = f(W_2 * h_1 + b_2)$, $y_{\text{hat}} = \text{Softmax}(W_3 * h_2 + b_3)$

Softmax Probability: $P(y=k|x) = \exp(z_k) / \sum_j \exp(z_j)$

Optimization & Loss Function

Cross-Entropy Loss: $L = - \sum y_k * \log(y_{\text{hat}_k}) + (\alpha / 2n) * \sum ||W_l||_F^2$

Training uses Backpropagation + Gradient Descent (Adam, L-BFGS, or SGD).

Meaning Breakdown & Mental Flowchart

input layer

↓

Receive raw input feature vector x

hidden layers

↓

Extract non-linear hierarchical feature embeddings

softmax output

↓

Produce normalized probability distribution across classes

backprop

↓

Propagate gradient error backwards to update weight matrices W_l

4. Perceptron

Mental Model

A single linear neuron that tests predictions and shifts its boundary weights only when it makes a mistake.

Mathematical Formulation & Prediction

Prediction: $y_{\text{hat}} = \text{sgn}(w^T * x + b)$

Condition to update: If $y_i * (w^T * x_i + b) \leq 0$ (Misclassified)

Optimization & Loss Function

Update Rule:

$w \leftarrow w + \eta * y_i * x_i$

$b \leftarrow b + \eta * y_i$

Meaning Breakdown & Mental Flowchart

correct prediction



Do nothing (weights stay unchanged)

misclassified sample



Shift weight vector w in the direction of the error $y_i * x_i$

simple loop



Iterate until linear convergence or max iterations reached

5. LogisticRegression

Mental Model

Compute a linear score for a sample, then map it through a Sigmoid curve to output a clean class probability between 0 and 1.

Mathematical Formulation & Prediction

Linear Score: $z = w^T * x + b$

Sigmoid Function: $P(y=1|x) = \sigma(z) = 1 / (1 + \exp(-z))$

Optimization & Loss Function

Binary Cross-Entropy Loss:

$$L(w) = - (1/n) * \sum [y_i * \log(p_i) + (1 - y_i) * \log(1 - p_i)] + (1 / 2C) * ||w||^2$$

Optimized via L-BFGS, Newton-CG, or Liblinear.

Meaning Breakdown & Mental Flowchart

linear score z



Combine features linearly $z = w^T * x + b$

sigmoid $\sigma(z)$



Squash z to probability scale $[0, 1]$

cross-entropy



Penalize confident wrong probability guesses exponentially

gradient descent



Adjust w to maximize log-likelihood of dataset

6. LogisticRegressionCV

Mental Model

Fit Logistic Regression across multiple internal cross-validation folds to automatically discover the best regularization strength C .

Mathematical Formulation & Prediction

$P(y=1|x) = \sigma(w^T * x + b)$

Validation Score: $CV_Score(C_k) = (1/K) * \sum_{v=1}^K Accuracy_val(C_k, Fold_v)$

Optimization & Loss Function

Solves internal grid search over a grid of C values:

$$C^{**} = \operatorname{argmax}_{\{C_k\}} \text{CV_Score}(C_k)$$

Final fit uses C^{**} on the full dataset.

Meaning Breakdown & Mental Flowchart

C grid search

↓

Test multiple regularization candidate values C

K-fold split

↓

Evaluate out-of-fold cross-validation accuracy for each C

select optimal C^*

↓

Pick C^* with maximum validation score

final model

↓

Fit Logistic Regression using optimal C^*

7. SVC (Support Vector Classifier)

Mental Model

Project data into higher-dimensional space using kernel functions where classes become linearly separable by a maximum-margin hyperplane.

Mathematical Formulation & Prediction

Dual Decision Function: $f(x) = \operatorname{sgn}(\sum_{i \in \text{SV}} \alpha_i * y_i * K(x_i, x) + b)$

Kernels $K(x_i, x)$:

- Linear: $x_i^T * x$

- RBF: $\exp(-\gamma * ||x_i - x||^2)$

- Poly: $(\gamma * x_i^T * x + r)^d$

Optimization & Loss Function

Dual Quadratic Optimization:

$$\max_{\alpha} \sum \alpha_i - 0.5 * \sum_i \sum_j \alpha_i * \alpha_j * y_i * y_j * K(x_i, x_j) \text{ s.t. } 0 \leq \alpha_i \leq C, \sum \alpha_i * y_i = 0$$

Solved using Sequential Minimal Optimization (SMO).

Meaning Breakdown & Mental Flowchart

kernel $K(x, x_i)$

↓

Compute non-linear similarity between samples in feature space

support vectors

↓

Identify key boundary points where $\alpha_i > 0$

max margin boundary

↓

Construct optimal separating boundary between classes

8. CalibratedClassifierCV

Mental Model

Wrap an uncalibrated classifier (like SVM or Naive Bayes) and pass its scores through a secondary mapping to output accurate, reliable probabilities.

Mathematical Formulation & Prediction

Platt Scaling (Sigmoid): $P(y=1|f) = 1 / (1 + \exp(A * f(x) + B))$

Isotonic Regression: $P(y=1|f) = g(f(x))$ where g is non-decreasing step function

Optimization & Loss Function

Fits parameters A, B via Maximum Likelihood on validation fold scores $f(x)$, or fits isotonic step curve g to minimize squared probability error.

Meaning Breakdown & Mental Flowchart

base model $f(x)$

↓

Get uncalibrated confidence score / distance $f(x)$

calibration map

↓

Pass $f(x)$ through Sigmoid curve (Platt) or Monotonic step (Isotonic)

calibrated $P(y|x)$

↓

Output true calibrated probability reflecting real-world frequencies

9. PassiveAggressiveClassifier

Mental Model

Be passive (do nothing) if a sample is correctly classified with margin; become aggressive (make huge step update) if error occurs.

Mathematical Formulation & Prediction

Hinge Loss: $L_t = \max(0, 1 - y_t * (w_t^T * x_t))$

Weight Update: $w_{t+1} = w_t + \tau_t * y_t * x_t$

Optimization & Loss Function

Step Size Tau: $\tau_t = \min(C, L_t / ||x_t||^2)$

Optimizes constrained distance minimization $||w - w_t||^2$ subject to linear loss constraint.

Meaning Breakdown & Mental Flowchart

sample correct ($L_t=0$)

↓

Passive mode: Keep weights $w_{t+1} = w_t$

sample error ($L_t>0$)

↓

Aggressive mode: Update weights proportional to loss L_t

C bound

↓

Clip step size τ_t by C to prevent noisy over-corrections

10. LabelPropagation

Mental Model

Build a similarity graph of all data points and let known class labels flow along graph edges to label unlabeled points.

Mathematical Formulation & Prediction

Graph Weight: $W_{ij} = \exp(-\gamma * ||x_i - x_j||^2)$

Label Transition: $Y_{t+1} = D^{-1} * W * Y_t$

Optimization & Loss Function

Iteratively multiplies probability label matrix Y by normalized transition matrix $D^{-1}W$, then clamps (resets) known labeled nodes back to their true values.

Meaning Breakdown & Mental Flowchart

similarity graph

↓

Connect data points with weights based on distance RBF kernel

label diffusion

↓

Propagate label distribution from labeled to unlabeled nodes

hard clamping

↓

Reset original ground-truth labels after each iteration until convergence

11. LabelSpreading

Mental Model

Propagate labels smoothly over a normalized graph Laplacian while allowing soft noise adjustments on original labels.

Mathematical Formulation & Prediction

Symmetric Normalized Graph: $S = D^{-1/2} * W * D^{-1/2}$

Iterative Update: $Y_{t+1} = \alpha * S * Y_t + (1 - \alpha) * Y_0$

Optimization & Loss Function

Closed-Form Solution:

$$Y^* = (1 - \alpha) * (I - \alpha * S)^{-1} * Y_0$$

α parameter balances graph smoothness versus initial label retention.

Meaning Breakdown & Mental Flowchart

normalized graph S

↓

Scale graph weights symmetrically to avoid degree bias

soft clamping α

↓

Blend graph consensus (α) with original labels ($1-\alpha$)

closed-form Y^*

↓

Solve exact steady-state label equilibrium matrix directly

12. RandomForestClassifier

Mental Model

Grow a forest of diverse decision trees on random data subsets and random feature splits, then combine their votes.

Mathematical Formulation & Prediction

Tree Vote: $P_m(y=c|x)$ from m -th tree trained on Bootstrap Sample B_m

Forest Output: $P(y=c|x) = (1/M) * \sum_{m=1}^M P_m(y=c|x)$

Optimization & Loss Function

Each tree node selects optimal split by sampling a random feature subset of size $\sqrt{D}$ and maximizing Gini Gain $\Delta I_G = I_G(\text{parent}) - [(N_L/N) I_G(L) + (N_R/N) I_G(R)]$.

Meaning Breakdown & Mental Flowchart

bootstrap sampling

↓

Train each tree on a random sample with replacement

feature randomization

↓

At each split node, search only a random subset of features

deep trees

↓

Grow unpruned trees to capture complex interactions

majority vote

↓

Average predictions to reduce variance drastically

13. GradientBoostingClassifier

Mental Model

Build decision trees sequentially, where each new tree focuses on correcting the specific errors (pseudo-residuals) of the previous ensemble.

Mathematical Formulation & Prediction

Ensemble Prediction: $F_m(x) = F_{m-1}(x) + \nu * \gamma_m * h_m(x)$

Pseudo-Residual: $r_{i,m} = - [\partial L(y_i, F(x_i)) / \partial F(x_i)]_{F=F_{m-1}}$

Optimization & Loss Function

Fits tree $h_m(x)$ to negative loss gradient $r_{i,m}$ using log-loss deviance, then scales by learning rate ν .

Meaning Breakdown & Mental Flowchart

current model F_{m-1}

↓

Generate predictions for all training samples

compute residuals r_m

↓

Calculate negative gradient direction of loss function

fit new tree h_m

↓

Train tree to predict error residuals r_m

add with shrinkage ν

↓

Update $F_m(x) = F_{m-1}(x) + \nu * h_m(x)$ to prevent overfitting

14. QuadraticDiscriminantAnalysis

Mental Model

Model each class as a separate multivariate Gaussian distribution with its own unique covariance matrix, creating quadratic decision boundaries.

Mathematical Formulation & Prediction

Discriminant Score:

$$\delta_k(x) = -0.5 * \log|\Sigma_k| - 0.5 * (x - \mu_k)^T * \Sigma_k^{-1} * (x - \mu_k) + \log(P(Y=k))$$

Optimization & Loss Function

Maximum Likelihood Estimation:

Estimates class means μ_k , class covariance matrices Σ_k , and prior probabilities $P(Y=k)$ directly from data.

Meaning Breakdown & Mental Flowchart

class mean μ_k

↓

Find average center location of class k

covariance Σ_k

↓

Fit class-specific shape, spread, and orientation matrix

quadratic boundary

↓

Compare quadratic distance scores and pick max $\delta_k(x)$

15. HistGradientBoostingClassifier

Mental Model

Discretize continuous features into integer bins to compute gradient histograms instantly, speeding up gradient boosting exponentially.

Mathematical Formulation & Prediction

Binning Transformation: $x_{\text{binned}} = \text{QuantileBin}(x, \text{max_bins}=255)$

Histogram Gradient: $G_{\text{bin}} = \sum_{i \in \text{Bin}} g_i$, $H_{\text{bin}} = \sum_{i \in \text{Bin}} h_i$

Optimization & Loss Function

Evaluates split thresholds over 256 discrete bins instead of sorting all individual sample values, reducing complexity from $O(N \log N)$ to $O(\text{max_bins})$.

Meaning Breakdown & Mental Flowchart

255 quantile bins

↓

Group continuous numbers into 256 integer buckets

gradient histogram

↓

Sum first (g_i) and second (h_i) derivatives per bin

instant split search

↓

Scan 256 bin boundaries to find best split threshold

super fast boosting



Train gradient boosted trees 10x-100x faster

16. RidgeClassifierCV

Mental Model

Cast classification as a multi-target linear regression with L2 regularization and tune parameter α via fast analytical Leave-One-Out Cross-Validation.

Mathematical Formulation & Prediction

Target Encoding: $Y \in \{-1, +1\}^{N \times C}$

Prediction: $y_{\text{hat}} = \text{argmax}_c (w_c^T x + b_c)$

Optimization & Loss Function

Analytical LOOCV Error:

$$\alpha^* = \text{argmin}_{\{\alpha_k\}} (1/n) * \sum [(y_i - y_{\text{hat}_{-i}}(\alpha_k))^2]$$

Computed instantly via SVD matrix decomposition without retraining.

Meaning Breakdown & Mental Flowchart

label matrix Y



Convert labels to -1/+1 multi-column indicators

SVD decomposition



Decompose feature matrix $X = U \Sigma V^T$ once

instant LOOCV score



Evaluate validation error for all α values simultaneously

select best α^*



Pick optimal α^* and compute final classifier weights

17. RidgeClassifier

Mental Model

Convert classification target labels to binary indicator numbers, fit a Ridge regression model, and pick the class with the highest score.

Mathematical Formulation & Prediction

Indicator Matrix: $Y \in \{-1, +1\}^{N \times C}$

Weight Solution: $w = (X^T X + \alpha * I)^{-1} * X^T * Y$

Optimization & Loss Function

Prediction Rule:

$$y_{\text{hat}} = \text{argmax}_c (w_c^T x + b_c)$$

Minimizes squared indicator error with L2 penalty.

Meaning Breakdown & Mental Flowchart

convert targets



Map class labels to -1/+1 columns

closed-form matrix

↓

Solve $w = (X^T X + \alpha I)^{-1} X^T Y$ in one step

argmax decision

↓

Output class corresponding to highest continuous linear output

18. AdaBoostClassifier

Mental Model

Combine weak decision stumps sequentially by boosting weights of misclassified samples after each round.

Mathematical Formulation & Prediction

Ensemble Output: $H(x) = \text{sgn}(\sum_{m=1}^M \alpha_m * h_m(x))$

Learner Weight: $\alpha_m = 0.5 * \log((1 - e_m) / e_m)$

Optimization & Loss Function

Sample Weight Update:

$w_{i, m+1} = w_{i, m} * \exp(-\alpha_m * y_i * h_m(x_i))$, normalized so $\sum w_i = 1$.

Implicitly minimizes Exponential Loss $L(y, f(x)) = \exp(-y * f(x))$.

Meaning Breakdown & Mental Flowchart

train weak stump

↓

Fit simple 1-split tree h_m on weighted data

calculate error e_m

↓

Compute weighted error rate of weak learner

compute α_m

↓

Assign voting weight α_m to current stump based on accuracy

re-weight samples

↓

Increase weights of misclassified data points for next stump

19. ExtraTreesClassifier

Mental Model

Build a forest of trees using random split thresholds for candidate features to maximize randomness and lower model variance.

Mathematical Formulation & Prediction

Split Threshold: $t_j \sim \text{Uniform}(\min(x_j), \max(x_j))$ across random feature subset

Optimization & Loss Function

Computes impurity gain for randomly drawn thresholds instead of searching for the exact mathematical optimum split point, reducing variance and computation time.

Meaning Breakdown & Mental Flowchart

random features



Select random subset of features at each node

random thresholds



Draw random split threshold t_j uniformly within range

pick best random split



Choose the best among the random candidates

forest voting



Average votes across trees to lower prediction variance

20. KNeighborsClassifier

Mental Model

Store all training data and classify a new point based on the majority vote of its K closest neighbors in feature space.

Mathematical Formulation & Prediction

Distance Metric: $d(\mathbf{x}, \mathbf{z}) = (\sum |\mathbf{x}_j - \mathbf{z}_j|^p)^{1/p}$

Class Probability: $P(Y=c|\mathbf{x}) = (1/K) * \sum_{\{i \in N_K(\mathbf{x}) \mid y_i = c\}}$

Optimization & Loss Function

Lazy Learning (No explicit training phase).

Search Acceleration: Uses KDTree or BallTree index structures to find nearest points fast.

Meaning Breakdown & Mental Flowchart

store samples



Keep training vectors in memory (or tree index)

find K neighbors



Compute distance to all points and select K closest

vote / average



Count class labels among K neighbors and predict majority class

21. BaggingClassifier

Mental Model

Train multiple copies of a base model on random bootstrap subsets of data and average their predictions to reduce model variance.

Mathematical Formulation & Prediction

Bootstrap Samples: $B_1, B_2, \dots, B_M$ drawn with replacement from dataset X

Ensemble Output: $f_{\text{bag}}(\mathbf{x}) = \text{Mode}(h_1(\mathbf{x}), h_2(\mathbf{x}), \dots, h_M(\mathbf{x}))$

Optimization & Loss Function

Fits base estimator h_m independently on each bootstrap subset B_m .

Meaning Breakdown & Mental Flowchart

create M subsets



Draw M bootstrap random samples with replacement

fit M base models



Train independent base model on each subset

combine predictions



Take majority vote / average probability across models

22. BernoulliNB

Mental Model

Calculate class probabilities assuming binary boolean features (0 or 1) that are independent given the class label.

Mathematical Formulation & Prediction

Likelihood: $P(\mathbf{x}|\mathbf{Y}=\mathbf{c}) = \prod_{j=1}^D [p_{\{c,j\}}^{x_j} * (1 - p_{\{c,j\}})^{(1 - x_j)}]$

Probability Parameter: $p_{\{c,j\}} = (N_{\{c,j\}} + \alpha) / (N_c + 2\alpha)$

Optimization & Loss Function

Decision Rule: $\hat{y} = \operatorname{argmax}_c [\log P(\mathbf{Y}=\mathbf{c}) + \sum \log P(x_j|\mathbf{Y}=\mathbf{c})]$

Uses Laplace smoothing α to prevent zero probability issues.

Meaning Breakdown & Mental Flowchart

binarize features



Check feature presence ($x_j=1$) or absence ($x_j=0$)

count frequencies



Calculate class probability $p_{\{c,j\}}$ for each feature

naive bayes rule



Multiply probabilities assuming conditional independence

23. LinearDiscriminantAnalysis

Mental Model

Project data onto a lower-dimensional axis that maximizes class separation relative to within-class variance, assuming shared covariance.

Mathematical Formulation & Prediction

Linear Discriminant Function:

$$\delta_k(\mathbf{x}) = \mathbf{x}^T \Sigma^{-1} * \mu_k - 0.5 * \mu_k^T \Sigma^{-1} * \mu_k + \log(P(\mathbf{Y}=\mathbf{k}))$$

Optimization & Loss Function

Assumes shared covariance matrix Σ across all classes, resulting in linear decision hyperplanes $\mathbf{w}_k^T * \mathbf{x} + b_k$.

Meaning Breakdown & Mental Flowchart

shared covariance Σ



Pool feature variances across all classes

calculate means μ_k

↓

Find average feature vector per class

linear decision surface

↓

Separate classes with linear hyperplane boundaries

24. GaussianNB

Mental Model

Assume continuous features within each class follow bell-shaped Gaussian distributions and combine their probabilities independently.

Mathematical Formulation & Prediction

Gaussian Likelihood:

$$P(\mathbf{x}_j | Y=c) = (1 / \sqrt{(2\pi * \sigma_{\{cj\}}^2)}) * \exp(- (\mathbf{x}_j - \mu_{\{cj\}})^2 / (2 * \sigma_{\{cj\}}^2))$$

Optimization & Loss Function

Estimates mean $\mu_{\{cj\}}$ and variance $\sigma_{\{cj\}}^2$ for each feature j per class c using Maximum Likelihood.

Meaning Breakdown & Mental Flowchart

fit gaussian curve

↓

Compute mean μ and variance σ^2 per feature per class

compute probability

↓

Evaluate Gaussian likelihood $P(\mathbf{x}_j|c)$ for input value

bayes rule decision

↓

Multiply feature likelihoods with class priors $P(c)$

25. NuSVC

Mental Model

Control SVM margin errors directly using parameter ν (nu) which bounds the fraction of support vectors and margin violations.

Mathematical Formulation & Prediction

Primal Optimization:

$$\min_{\{\mathbf{w}, \mathbf{b}, \xi, \rho\}} 0.5 * \|\mathbf{w}\|^2 - \nu * \rho + (1/n) * \sum \xi_i \text{ s.t. } \mathbf{y}_i * (\mathbf{w}^T \phi(\mathbf{x}_i) + \mathbf{b}) \geq \rho - \xi_i$$

Optimization & Loss Function

Parameter $\nu \in (0, 1]$ automatically sets an upper bound on margin error fraction and lower bound on support vector count.

Meaning Breakdown & Mental Flowchart

nu parameter ν

↓

Set target percentage of allowed margin errors

flexible margin ρ

↓

Dynamically adjust margin width parameter ρ

SMO solver

↓

Solve dual QP formulation via LIBSVM

26. DecisionTreeClassifier

Mental Model

Ask a series of binary yes/no questions on features to recursively split data into increasingly pure class regions.

Mathematical Formulation & Prediction

Node Impurity $I(A)$:

- Gini: $1 - \sum p_k^2$

- Entropy: $-\sum p_k \log_2(p_k)$

Information Gain: $\Delta I = I(\text{Parent}) - [(N_L/N) I(L) + (N_R/N) I(R)]$

Optimization & Loss Function

Finds feature j and split threshold t maximizing Information Gain ΔI recursively until stopping criteria are met.

Meaning Breakdown & Mental Flowchart

evaluate all splits

↓

Test feature thresholds to find maximum purity increase

split node

↓

Divide data into left and right child nodes

repeat recursively

↓

Continue splitting until leaf nodes are pure or max depth reached

27. NearestCentroid

Mental Model

Represent each class by its mean center vector (centroid) and assign new samples to the closest centroid.

Mathematical Formulation & Prediction

Class Centroid: $\mu_c = (1 / N_c) * \sum_{i: y_i = c} x_i$

Prediction: $y_{\text{hat}} = \operatorname{argmin}_c d(x, \mu_c)$

Optimization & Loss Function

Calculates Euclidean or Manhattan distance to each class mean vector μ_c .

Meaning Breakdown & Mental Flowchart

calculate centroids

↓

Compute average center point μ_c for each class

measure distance

↓

Calculate distance from sample x to all centroids

assign class

↓

Pick class label of the closest centroid

28. ExtraTreeClassifier

Mental Model

Build a single extremely randomized decision tree where candidate split thresholds are drawn completely at random.

Mathematical Formulation & Prediction

Split Threshold: $t_j \sim \text{Uniform}(\min(x_j), \max(x_j))$

Optimization & Loss Function

Selects random split thresholds for candidate features and picks the best among them, speeding up tree construction.

Meaning Breakdown & Mental Flowchart

draw random t_j

↓

Pick random split threshold for candidate features

evaluate impurity

↓

Compute Gini impurity for random threshold splits

split & repeat

↓

Split data fast with lower computational overhead

29. CheckingClassifier

Mental Model

An internal diagnostic test model used to check Scikit-Learn estimator API compliance and data input formatting.

Mathematical Formulation & Prediction

Internal API Check: Validate X shape, y format, and fit/predict method signatures.

Optimization & Loss Function

Performs no real ML training; verifies estimator contract during automated testing.

Meaning Breakdown & Mental Flowchart

receive input X, y

↓

Accept data arrays

validate contract

↓

Check API compatibility and parameter constraints

return mock output

↓

Provide mock test response

30. DummyClassifier

Mental Model

Make simple non-smart baseline predictions (like always predicting the majority class) to benchmark real ML model uplift.

Mathematical Formulation & Prediction

Strategy 'most_frequent': $P(y=c|x) = 1$ if $c = \text{Mode}(Y)$ else 0

Strategy 'stratified': $P(y=c|x) = \text{Prior}(c) = N_c / N$

Optimization & Loss Function

Computes simple summary statistics on target labels Y without looking at input features X.

Meaning Breakdown & Mental Flowchart

ignore features X

↓

Do not inspect input feature values

calculate label stat

↓

Count class frequencies in training targets Y

predict baseline

↓

Output majority class or random prior distribution

Chapter 3: Regression Models & Architectures (36 Models)

31. SVR

Mental Model

Fit a prediction tube around data points where errors within threshold ε are completely ignored, penalizing only points outside.

Mathematical Formulation & Prediction

Dual Function: $f(x) = \sum_{i \in \text{SV}} (\alpha_i - \alpha_i^*) * K(x_i, x) + b$

ε -Insensitive Loss: $L_\varepsilon(y, f(x)) = \max(0, |y - f(x)| - \varepsilon)$

Optimization & Loss Function

$\min_{\{w, b, \xi, \xi^*\}} 0.5 * ||w||^2 + C * \sum (\xi_i + \xi_i^*)$
s.t. $|y_i - (w^T \phi(x_i) + b)| \leq \varepsilon + \xi_i$

Meaning Breakdown & Mental Flowchart

insensitive tube ε

↓

Ignore small errors falling within distance ε of tube

slack variables ξ_i

↓

Penalize points falling outside the ε -tube linearly

C parameter

↓

Control trade-off between tube smoothness and error penalty

32. RandomForestRegressor

Mental Model

Grow multiple regression trees on random data samples and average their continuous predictions to reduce model variance.

Mathematical Formulation & Prediction

Tree Output: $T_m(x)$ for m -th regression tree

Ensemble Output: $f(x) = (1/M) * \sum_{m=1}^M T_m(x)$

Optimization & Loss Function

Node Split Criterion: Minimizes Variance / MSE

Split Loss = $\sum_{i \in \text{Left}} (y_i - c_1)^2 + \sum_{i \in \text{Right}} (y_i - c_2)^2$

Meaning Breakdown & Mental Flowchart

bootstrap samples

↓

Train trees on independent random data subsets

feature subsets

↓

Select random feature candidate subsets at each split node

average predictions

↓

Average tree outputs to produce smooth continuous predictions

33. ExtraTreesRegressor

Mental Model

Average predictions across regression trees built with completely random split thresholds for candidate features.

Mathematical Formulation & Prediction

Split Threshold: $t_j \sim \text{Uniform}(\min(x_j), \max(x_j))$

Ensemble Prediction: $f(x) = (1/M) * \sum_{m=1}^M T_m(x)$

Optimization & Loss Function

Evaluates random split thresholds to reduce tree correlation and overall prediction variance.

Meaning Breakdown & Mental Flowchart

random thresholds

↓

Draw random split threshold t_j per candidate feature

fit extra trees

↓

Build highly randomized regression trees

average ensemble

↓

Average predictions for robust, low-variance regression

34. AdaBoostRegressor

Mental Model

Train sequence of weak regression trees, increasing weights of samples with large relative error after each stage.

Mathematical Formulation & Prediction

Relative Error: $e_{im} = |y_i - T_m(x_i)| / \text{Max_Error}$

Ensemble Prediction: Weighted Median of $\{T_m(x)\}$ using weights $\log(1/\beta_m)$

Optimization & Loss Function

Updates sample weights $w_{i, m+1} = w_{i, m} * \beta_m^{\{(1 - e_{im})\}}$ where $\beta_m = e_m / (1 - e_m)$.

Meaning Breakdown & Mental Flowchart

fit weak tree

↓

Train regression tree T_m on weighted sample data

calculate errors

↓

Compute relative error e_{im} for each sample

weighted median

↓

Combine tree outputs using weighted median prediction

35. NuSVR

Mental Model

Automatically control the width of the insensitive error tube ε using parameter ν (nu).

Mathematical Formulation & Prediction

Objective:

$$\min_{\{w, b, \xi, \xi^*, \varepsilon\}} 0.5 * ||w||^2 + C * (\nu * \varepsilon + (1/n) * \sum (\xi_i + \xi_i^*))$$

$$\text{s.t. } |y_i - (w^T \phi(x_i) + b)| \leq \varepsilon + \xi_i$$

Optimization & Loss Function

Parameter $\nu \in (0, 1]$ bounds the fraction of support vectors and automatically determines tube width ε .

Meaning Breakdown & Mental Flowchart

nu parameter ν



Set desired fraction of support vector data points

auto tube width ε



Automatically compute optimal insensitive tube margin ε

SMO optimization



Solve kernelized dual SVR problem via LIBSVM

36. GradientBoostingRegressor

Mental Model

Fit decision trees sequentially to the residual errors (gradients) of the existing ensemble.

Mathematical Formulation & Prediction

$$\text{Ensemble Update: } F_m(x) = F_{m-1}(x) + \nu * \gamma_m * h_m(x)$$

$$\text{Gradient Residual: } r_{im} = - [\partial L(y_i, F(x_i)) / \partial F(x_i)]_{F=F_{m-1}}$$

Optimization & Loss Function

Loss Functions: 'squared_error' (MSE), 'absolute_error' (MAE), 'huber', 'quantile'.

Meaning Breakdown & Mental Flowchart

calculate residuals



Compute error $r_{im} = y_i - F_{m-1}(x_i)$

fit new tree h_m



Train tree h_m to predict residuals r_{im}

apply shrinkage ν



Scale tree output by learning rate ν and add to ensemble

37. KNeighborsRegressor

Mental Model

Predict a continuous target by averaging the target values of the K nearest neighbors in feature space.

Mathematical Formulation & Prediction

$$\text{Prediction: } f(x) = \sum_{i \in N_K(x)} w_i * y_i / \sum w_i$$

$$w_i = 1 \text{ (Uniform) or } w_i = 1 / d(x, x_i) \text{ (Distance-weighted)}$$

Optimization & Loss Function

Uses Minkowski distance $d(x, z) = (\sum |x_j - z_j|^p)^{1/p}$ with KDTree/BallTree search.

Meaning Breakdown & Mental Flowchart

find K neighbors

↓

Locate K nearest training samples by distance

weight targets

↓

Assign weights w_i based on inverse distance

weighted average

↓

Compute weighted target mean to produce prediction $f(x)$

38. HistGradientBoostingRegressor

Mental Model

Discretize continuous inputs into integer bins to compute gradient histograms rapidly for fast regression tree splits.

Mathematical Formulation & Prediction

Binning: $x_{\text{binned}} = \text{QuantileBin}(x, \text{max_bins}=255)$

Tree Split: Finds bin boundary maximizing loss reduction using gradient/hessian sums per bin.

Optimization & Loss Function

Accelerates gradient boosting on large datasets while natively supporting missing values.

Meaning Breakdown & Mental Flowchart

bin features

↓

Group continuous feature values into 256 discrete bins

sum gradients

↓

Accumulate gradient & hessian sums per bin

rapid split search

↓

Scan 256 bin boundaries for optimal split threshold

39. BaggingRegressor

Mental Model

Average continuous outputs from multiple base regressors trained on random bootstrap samples.

Mathematical Formulation & Prediction

Bootstrap Samples: $B_1, B_2, \dots, B_M$

Prediction: $f_{\text{bag}}(x) = (1/M) * \sum_{m=1}^M h_m(x)$

Optimization & Loss Function

Fits base regressor h_m independently on each random bootstrap data subset.

Meaning Breakdown & Mental Flowchart

create M samples



Draw bootstrap random data subsets with replacement

fit regressors



Train independent base regressor on each subset

average outputs



Average predicted values across all M base regressors

40. MLPRegressor

Mental Model

Pass features through hidden neural layers and output a continuous prediction via a linear output neuron.

Mathematical Formulation & Prediction

Hidden Layers: $h_1 = f(W_1 * x + b_1)$, $h_2 = f(W_2 * h_1 + b_2)$

Linear Output: $f(x) = W_3 * h_2 + b_3$

Optimization & Loss Function

MSE Loss: $L = 0.5 * (1/n) * \sum (y_i - f(x_i))^2 + (\alpha / 2n) * \sum ||W_l||_F^2$

Optimized via Backpropagation with Adam or L-BFGS.

Meaning Breakdown & Mental Flowchart

hidden activations



Compute non-linear feature representations in hidden layers

linear output neuron



Produce unconstrained continuous target value $f(x)$

MSE loss backprop



Compute squared error and backpropagate weight updates

41. HuberRegressor

Mental Model

Fit a robust linear model that uses squared loss for small errors and linear loss for large outlier errors.

Mathematical Formulation & Prediction

Huber Loss Function $H_\delta(z)$:

- $0.5 * z^2$ if $|z| \leq \delta$

- $\delta * |z| - 0.5 * \delta^2$ if $|z| > \delta$

Optimization & Loss Function

$\min_{\{w, \sigma\}} \sum [\sigma + H_\delta((y_i - w^T x_i) / \sigma) * \sigma] + \alpha * ||w||_2^2$

Optimizes weights w and scale parameter σ simultaneously.

Meaning Breakdown & Mental Flowchart

small errors ($|z| \leq \delta$)

↓

Behave like MSE (quadratic smooth loss)

large errors ($|z| > \delta$)

↓

Behave like MAE (linear loss; immune to outlier blowup)

scale parameter σ

↓

Scale residuals dynamically to identify outlier threshold

42. LinearSVR

Mental Model

Fit a linear prediction function using ε -insensitive loss solved rapidly via LIBLINEAR.

Mathematical Formulation & Prediction

Prediction: $f(x) = \hat{w}^T x + b$

Loss: $L_\varepsilon = \max(0, |y_i - (\hat{w}^T x_i + b)| - \varepsilon)$

Optimization & Loss Function

$\min_{\{w,b\}} 0.5 * ||w||^2 + C * \sum \max(0, |y_i - (\hat{w}^T x_i + b)| - \varepsilon)$

Fast linear solver using Coordinate Descent.

Meaning Breakdown & Mental Flowchart

linear prediction

↓

Calculate linear prediction $\hat{w}^T x + b$

ε -insensitive loss

↓

Ignore errors inside ε -tube margin

LIBLINEAR solver

↓

Optimize linear weights fast for large datasets

43. RidgeCV

Mental Model

Fit Ridge regression and select the optimal L2 penalty α using efficient analytical Leave-One-Out Cross-Validation.

Mathematical Formulation & Prediction

Analytical LOOCV MSE:

LOOCV_Error(α_k) = $(1/n) * \sum [(y_i - \hat{y}_{-i}(\alpha_k))^2]$

Optimization & Loss Function

Computes LOOCV error for all α values simultaneously using SVD decomposition without refitting models.

Meaning Breakdown & Mental Flowchart

SVD decomposition

↓

Decompose feature matrix X once via SVD

analytical LOOCV

↓

Compute leave-one-out validation error for all α candidates

select optimal α^*

↓

Pick α^* with minimal validation MSE and return final weights

44. BayesianRidge

Mental Model

Estimate linear regression parameters probabilistically under Gaussian priors, providing prediction uncertainty.

Mathematical Formulation & Prediction

Posterior Distribution: $P(\mathbf{w}|\mathbf{y}, \mathbf{X}) \sim \mathcal{N}(\boldsymbol{\mu}_N, \boldsymbol{\Sigma}_N)$

$$\boldsymbol{\mu}_N = \alpha^* \boldsymbol{\Sigma}_N^* \mathbf{X}^T \mathbf{y}, \boldsymbol{\Sigma}_N = (\lambda^* \mathbf{I} + \alpha^* \mathbf{X}^T \mathbf{X})^{-1}$$

Optimization & Loss Function

Estimates precision hyperparameters α (noise precision) and λ (weight prior precision) via Empirical Bayes (Type-II ML).

Meaning Breakdown & Mental Flowchart

gaussian prior

↓

Assume Gaussian weight prior $P(\mathbf{w}|\lambda) = \mathcal{N}(0, \lambda^{-1} \mathbf{I})$

posterior $\boldsymbol{\mu}_N, \boldsymbol{\Sigma}_N$

↓

Calculate exact posterior weight mean and covariance

prediction variance

↓

Output expected prediction mean and uncertainty variance

45. Ridge

Mental Model

Shrink linear regression coefficients towards zero using an L2 penalty to prevent overfitting and handle multicollinearity.

Mathematical Formulation & Prediction

Objective: $\min_{\mathbf{w}} \|\mathbf{y} - \mathbf{X} \mathbf{w}\|_2^2 + \alpha^* \|\mathbf{w}\|_2^2$

Closed-Form Solution: $\mathbf{w}^* = (\mathbf{X}^T \mathbf{X} + \alpha^* \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$

Optimization & Loss Function

Adds L2 penalty term $\alpha^* \|\mathbf{w}\|^2$ to RSS, making $(\mathbf{X}^T \mathbf{X} + \alpha^* \mathbf{I})$ invertible even when features are correlated.

Meaning Breakdown & Mental Flowchart

residual sum squares

↓

Minimize prediction error $\|\mathbf{y} - \mathbf{X} \mathbf{w}\|^2$

L2 penalty $\alpha^* \|\mathbf{w}\|^2$

↓

Shrink coefficient magnitudes towards zero

closed-form matrix

↓

Solve $w^* = (X^T X + \alpha I)^{-1} X^T y$ in one step

46. LinearRegression

Mental Model

Find the optimal linear hyperplane that minimizes the sum of squared distances to all training data points.

Mathematical Formulation & Prediction

Prediction: $y_{\text{hat}} = w^T x + b$

Objective: $\min_{\{w\}} ||y - X w||^2_2$

Optimization & Loss Function

Closed-Form Ordinary Least Squares (OLS):

$$w^* = (X^T X)^{-1} X^T y$$

Computed via Singular Value Decomposition (SVD) of X .

Meaning Breakdown & Mental Flowchart

linear combination

↓

Calculate $y_{\text{hat}} = w_1 x_1 + \dots + w_d x_d + b$

squared residuals

↓

Compute sum of squared errors $\sum (y_i - y_{\text{hat}_i})^2$

SVD OLS solver

↓

Find global minimum weights analytically via SVD

47. TransformedTargetRegressor

Mental Model

Transform non-linear target y (e.g., log transformation), train a regressor, and inverse-transform predictions back to real scale.

Mathematical Formulation & Prediction

Forward Target Transform: $y_{\text{trans}} = g(y)$

Fit: Regressor trained on (X, y_{trans})

Inverse Transform: $y_{\text{pred}} = g^{-1}(f(x))$

Optimization & Loss Function

Applies forward mapping g during `fit()` and inverse mapping g^{-1} during `predict()`.

Meaning Breakdown & Mental Flowchart

transform $y \rightarrow g(y)$

↓

Apply log / sqrt / custom transform to target y

fit base regressor

↓

Train underlying regressor on transformed target

inverse $g^{-1}(y_{\text{hat}})$

↓

Apply inverse transformation to return predictions to real units

48. LassoCV

Mental Model

Compute the full Lasso L1 regularization path across an α grid and pick the best α using K-fold Cross-Validation.

Mathematical Formulation & Prediction

Lasso Path: $w(\alpha)$ computed via **Coordinate Descent** for $\alpha_1 > \alpha_2 > \dots > \alpha_K$

Validation Score: $\alpha^* = \operatorname{argmin}_{\{\alpha_k\}} (1/K) * \sum_{v=1}^K \text{MSE}_{\{\text{val}\}}(\alpha_k, \text{Fold}_v)$

Optimization & Loss Function

Computes entire Lasso regularization path efficiently along an α grid using warm-restarted Coordinate Descent.

Meaning Breakdown & Mental Flowchart

α **grid path**

↓

Generate grid of regularization parameter values α

coordinate descent

↓

Compute sparse Lasso coefficient path $w(\alpha)$

select optimal α^*

↓

Pick α^* with minimal cross-validation MSE

49. ElasticNetCV

Mental Model

Cross-validate over a 2D grid of penalty scale α and L1 ratio to find the optimal mix of Lasso and Ridge regularization.

Mathematical Formulation & Prediction

Objective: $\min_{\{w\}} (1/2n) \|y - Xw\|^2 + \alpha * [l1_ratio * \|w\|_1 + 0.5 * (1 - l1_ratio) * \|w\|_2^2]$

CV Selection: $(\alpha^*, l1_ratio^*) = \operatorname{argmin} \text{CV_MSE}(\alpha, l1_ratio)$

Optimization & Loss Function

Performs grid search over combinations of α and $l1_ratio$ via Coordinate Descent.

Meaning Breakdown & Mental Flowchart

2D parameter grid

↓

Set candidate values for α and $l1_ratio$

coordinate descent

↓

Fit ElasticNet for each parameter pair across CV folds

optimal pair

↓

Select optimal $(\alpha^*, l1_ratio^*)$ with minimum validation error

50. LassoLarsCV

Mental Model

Compute the exact piecewise linear Lasso path using Least Angle Regression (LARS) and select α via Cross-Validation.

Mathematical Formulation & Prediction

LARS Lasso Path: Computes exact piecewise linear breakpoints where features enter/exit the model.

Optimization & Loss Function

Evaluates cross-validation error along exact LARS breakpoints rather than a fixed heuristic grid.

Meaning Breakdown & Mental Flowchart

LARS path

↓

Compute exact piecewise linear Lasso coefficient path

cross-validation

↓

Evaluate prediction error at each path breakpoint

select best step

↓

Pick optimal regularization step along LARS path

51. LassoLarsIC

Mental Model

Select the optimal Lasso parameter α along the LARS path using Information Criteria (AIC or BIC) without cross-validation folds.

Mathematical Formulation & Prediction

$$\text{AIC} = n * \log(\text{MSE}) + 2 * k$$

$$\text{BIC} = n * \log(\text{MSE}) + k * \log(n)$$

where $k = ||w||_0$ (number of non-zero active features)

Optimization & Loss Function

Selects optimal step along LARS path in a single training pass using analytical AIC or BIC score.

Meaning Breakdown & Mental Flowchart

LARS path

↓

Generate exact piecewise linear feature entry path

compute AIC / BIC

↓

Evaluate information criterion score for each path step

instant selection

↓

Select optimal model without performing K-fold splits

52. LarsCV

Mental Model

Fit Least Angle Regression (LARS) and select the optimal number of steps using Cross-Validation.

Mathematical Formulation & Prediction

LARS Feature Direction: Move weights w in direction equiangular to features with highest correlation $c_j = x_j^T (y - \hat{y})$.

Optimization & Loss Function

Evaluates cross-validation error along the LARS forward path to select optimal step count.

Meaning Breakdown & Mental Flowchart

equiangular path

↓

Advance weights equiangularly as features enter active set

cross-validation

↓

Measure out-of-fold MSE along LARS steps

select step count

↓

Choose optimal number of features to retain

53. Lars

Mental Model

Stepwise forward feature selection algorithm that advances weight coefficients along equiangular correlation directions.

Mathematical Formulation & Prediction

Correlation Vector: $c_j = x_j^T (y - \hat{y}_{k-1})$

Equiangular Vector: Advance w in direction having equal correlation with all active features.

Optimization & Loss Function

Computes entire forward feature selection path in same order of computational complexity as a single OLS fit.

Meaning Breakdown & Mental Flowchart

start at $w=0$

↓

Initialize all weights to zero

find max correlation

↓

Select feature most correlated with current residual

move equiangularly

↓

Advance weights until another feature reaches equal correlation

54. SGDRegressor

Mental Model

Optimize linear regression models iteratively using Stochastic Gradient Descent updates on mini-batches.

Mathematical Formulation & Prediction

Gradient Update: $w_{t+1} = w_t - \eta_t * [(w_t^T x_i - y_i) * x_i + \alpha * \nabla R(w_t)]$

Optimization & Loss Function

Loss Options: 'squared_error' (MSE), 'huber', 'epsilon_insensitive'. Regularization: L2, L1, or ElasticNet.

Meaning Breakdown & Mental Flowchart

mini-batch sample

↓

Fetch random sample/batch from training data

calculate gradient

↓

Compute loss gradient ∇L with respect to weights w

update weights

↓

Apply SGD step $w_{t+1} = w_t - \eta * \nabla L$

55. RANSACRegressor

Mental Model

Fit a linear model on random subsets of data to separate true inlier points from extreme outlier noise.

Mathematical Formulation & Prediction

Inlier Condition: $|y_i - f_{\text{sub}}(x_i)| \leq t_{\text{threshold}}$

Optimization & Loss Function

Iteratively fits base regressor on random sample subsets, counts consensus inliers within threshold t , and refits final model on the largest inlier set.

Meaning Breakdown & Mental Flowchart

random sample

↓

Select minimum random data subset to fit trial model

count inliers

↓

Identify all data points falling within residual threshold t

largest consensus

↓

Find subset generating maximum inlier consensus set

refit final model

↓

Re-estimate final regressor weights using inliers only

56. ElasticNet

Mental Model

Combine L1 (Lasso) feature selection with L2 (Ridge) grouping penalty to handle correlated features while maintaining sparsity.

Mathematical Formulation & Prediction

Objective:

$$\min_{\{w\}} (1/2n) * ||y - X w||_2^2 + \alpha * l1_ratio * ||w||_1 + 0.5 * \alpha * (1 - l1_ratio) * ||w||_2^2$$

Optimization & Loss Function

Solved via Coordinate Descent using soft-thresholding updates.

Meaning Breakdown & Mental Flowchart

L1 penalty



Enforce sparsity by driving irrelevant feature weights to zero

L2 penalty



Encourage group selection among correlated feature variables

coordinate descent



Iteratively optimize each weight w_j while holding others fixed

57. Lasso

Mental Model

Apply an L1 absolute weight penalty to shrink unnecessary feature coefficients strictly to zero, performing automatic feature selection.

Mathematical Formulation & Prediction

Objective: $\min_{\mathbf{w}} \left\{ \frac{1}{2n} \|\mathbf{y} - \mathbf{X} \mathbf{w}\|_2^2 + \alpha \|\mathbf{w}\|_1 \right\}$

Soft-Thresholding Update: $\mathbf{w}_j = S(\mathbf{X}_j^T (\mathbf{y} - \mathbf{X}_{-j} \mathbf{w}_{-j}), \alpha) / \|\mathbf{X}_j\|^2$

Optimization & Loss Function

Solved via Coordinate Descent with Soft-Thresholding operator $S(z, \alpha) = \text{sgn}(z) * \max(0, |z| - \alpha)$.

Meaning Breakdown & Mental Flowchart

L1 penalty $\alpha \|\mathbf{w}\|_1$



Add sum of absolute weights to loss function

soft-thresholding



Pull small coefficients to exact zero $w_j = 0$

sparse feature set



Produce interpretable model with sparse non-zero coefficients

58. OrthogonalMatchingPursuitCV

Mental Model

Cross-validate over the number of non-zero features in Orthogonal Matching Pursuit to select optimal feature sparsity.

Mathematical Formulation & Prediction

OMP Path: Solves $\min \|\mathbf{y} - \mathbf{X} \mathbf{w}\|_2^2$ s.t. $\|\mathbf{w}\|_0 \leq k$ for $k = 1, 2, \dots$

CV Selection: $k^* = \text{argmin}_k \text{CrossValidation_MSE}(k)$

Optimization & Loss Function

Selects optimal sparsity parameter k using out-of-fold cross-validation error.

Meaning Breakdown & Mental Flowchart

OMP path



Generate greedy sparse feature selection path

cross-validation

↓

Evaluate out-of-fold prediction MSE for feature counts k

select sparsity k^*

↓

Choose optimal non-zero feature count k^*

59. ExtraTreeRegressor

Mental Model

Build a single extremely randomized regression tree using uniform random split thresholds.

Mathematical Formulation & Prediction

Split Threshold: $t_j \sim \text{Uniform}(\min(x_j), \max(x_j))$

Leaf Output: Average target value y_{mean} in leaf region

Optimization & Loss Function

Draws random thresholds per candidate feature to construct single fast randomized regression tree.

Meaning Breakdown & Mental Flowchart

random t_j

↓

Draw random split threshold for feature candidate

variance reduction

↓

Build fast randomized single regression tree

leaf mean

↓

Output average target value of samples in leaf node

60. PassiveAggressiveRegressor

Mental Model

Stay passive if prediction error is within threshold ε ; aggressively update weights if error exceeds ε .

Mathematical Formulation & Prediction

Loss: $L_{\varepsilon} = \max(0, |y_t - w_t^T x_t| - \varepsilon)$

Update: $w_{t+1} = w_t + \text{sgn}(y_t - w_t^T x_t) * \tau_t * x_t$

where $\tau_t = \min(C, L_{\varepsilon} / \|x_t\|^2)$

Optimization & Loss Function

Online margin-based continuous regression update rule.

Meaning Breakdown & Mental Flowchart

error within ε

↓

Passive mode: Keep current weights $w_{t+1} = w_t$

error exceeds ε

↓

Aggressive mode: Update weights proportional to error loss

C bound

↓

Clip step size τ_t by C to prevent noisy weight jumps

61. GaussianProcessRegressor

Mental Model

Model target values as a continuous Gaussian Process, providing non-parametric predictions with complete uncertainty bounds.

Mathematical Formulation & Prediction

Prior: $f(x) \sim \text{GP}(m(x), k(x, x'))$

Posterior Mean: $\mu^* = K(x^*, X) * [K(X, X) + \sigma_n^2 I]^{-1} * y$

Posterior Variance: $\Sigma^* = K(x^*, x^*) - K(x^*, X) * [K(X, X) + \sigma_n^2 I]^{-1} * K(X, x^*)$

Optimization & Loss Function

Maximizes Marginal Log-Likelihood to optimize kernel hyperparameters.

Meaning Breakdown & Mental Flowchart

kernel covariance K

↓

Measure similarity between points via kernel $k(x, x')$

posterior mean μ^*

↓

Compute expected continuous target prediction

confidence bounds

↓

Calculate variance Σ^* to output 95% uncertainty confidence interval

62. OrthogonalMatchingPursuit

Mental Model

Greedy forward feature selection algorithm that selects features most correlated with target residuals and projects orthogonally.

Mathematical Formulation & Prediction

Objective: $\min_{\{w\}} \|y - X w\|_2^2$ s.t. $\|w\|_0 \leq n_{\text{nonzero_coefs}}$

Residual Update: $r_k = y - X_{\{S_k\}} w_{\{S_k\}}$

Optimization & Loss Function

At each step, picks feature x_j most correlated with current residual $r_{\{k-1\}}$, adds it to active set S_k , and projects target orthogonally onto $X_{\{S_k\}}$.

Meaning Breakdown & Mental Flowchart

pick max correlated x_j

↓

Find feature most correlated with residual vector r

add to active set S

↓

Include feature x_j in active feature subset

orthogonal projection

↓

Project target y orthogonally onto selected features and update residual r

63. DecisionTreeRegressor

Mental Model

Recursively split feature space into rectangular regions and predict the mean target value of training samples in each leaf region.

Mathematical Formulation & Prediction

Node MSE: $MSE = (1/N) * \sum_{i \in \text{Region}} (y_i - y_{\text{mean}})^2$

Split Optimization: $\min_{[j, s]} [MSE(\text{Left}) + MSE(\text{Right})]$

Optimization & Loss Function

Finds feature j and threshold s minimizing Mean Squared Error (or MAE) recursively until leaf stopping criteria are reached.

Meaning Breakdown & Mental Flowchart

test feature splits

↓

Evaluate thresholds across features to minimize MSE

partition region

↓

Split feature space into left and right sub-regions

leaf mean prediction

↓

Predict average target y_{mean} for any sample landing in leaf

64. DummyRegressor

Mental Model

Make simple non-smart baseline predictions (like always predicting target mean or median) to benchmark real ML model performance.

Mathematical Formulation & Prediction

Strategy 'mean': $y_{\text{pred}} = \text{Mean}(Y)$

Strategy 'median': $y_{\text{pred}} = \text{Median}(Y)$

Strategy 'quantile': $y_{\text{pred}} = \text{Quantile}(Y, q)$

Optimization & Loss Function

Calculates simple summary statistic on training target values Y without inspecting input features X .

Meaning Breakdown & Mental Flowchart

ignore features X

↓

Do not inspect input feature columns

calculate target summary

↓

Compute mean or median of target column Y

output constant baseline



Return constant baseline prediction for all inputs

65. LassoLars

Mental Model

Solve Lasso L1 regularized linear regression along the exact piecewise linear Least Angle Regression (LARS) path.

Mathematical Formulation & Prediction

Objective: $\min_{\{w\}} (1/2n) * ||y - X w||^2_2 + \alpha * ||w||_1$

LARS Algorithm Variant: Modifies LARS to handle feature drops when coefficients hit zero.

Optimization & Loss Function

Computes exact Lasso regularization path in computationally efficient LARS steps.

Meaning Breakdown & Mental Flowchart

LARS path modification



Track features entering model and drop features hitting zero weight

piecewise linear $w(\alpha)$



Compute exact analytical coefficient trajectory

exact Lasso weights



Return sparse weights w for target penalty α

66. KernelRidge

Mental Model

Combine L2 regularized Ridge regression with non-linear kernel functions to fit smooth non-linear regression curves.

Mathematical Formulation & Prediction

Dual Objective: $\min_{\{\alpha\}} ||y - K \alpha||^2_2 + \gamma * \alpha^T * K * \alpha$

Closed-Form Solution: $\alpha^* = (K + \gamma * I)^{-1} * y$

Prediction: $f(x) = \sum_{i=1}^n \alpha_i * K(x_i, x)$

Optimization & Loss Function

Solves dual kernelized Ridge regression analytically via matrix inversion.

Meaning Breakdown & Mental Flowchart

kernel matrix K



Compute non-linear kernel matrix $K_{\{ij\}} = K(x_i, x_j)$

closed-form dual α^*



Solve $\alpha^* = (K + \gamma * I)^{-1} y$ in a single matrix inversion step

non-linear prediction



Evaluate smooth non-linear regression function $f(x) = \sum \alpha_i K(x_i, x)$

Chapter 4: Computer Vision & Deep Learning Mental Models

4.1 Classic CNNs & Vision Transformers

Model / Architecture	Category	Mental Model & Core Idea	Primary Application
AlexNet	Classic CNN	Early deep CNN with ReLU and Dropout that revolutionized ImageNet classification.	Historical baseline CNN
VGG	Classic CNN	Very deep stack of simple 3x3 convolution filters and max pooling layers.	Classic feature extraction
ResNet	Residual CNN	Introduced skip connections (identity shortcuts) so 100+ layer networks train smoothly.	Standard visual backbone
DenseNet	Dense CNN	Connects every layer to all subsequent layers to maximize feature reuse.	Feature-reuse classification
Inception	Multi-Scale CNN	Runs multi-scale 1x1, 3x3, and 5x5 convolution filters in parallel in one block.	Multi-scale visual patterns
EfficientNet	Compounded CNN	Systematically scales depth, width, and resolution using compound coefficient.	High accuracy per compute
ConvNeXt	Modern CNN	Modernized CNN architecture incorporating design principles from Vision Transformers.	Modern state-of-art CNN
ViT (Vision Transformer)	Vision Transformer	Splits images into 16x16 patches and processes them as tokens via self-attention.	Global visual relationships
DeiT	Data-Efficient ViT	Uses teacher-student distillation to train Transformers efficiently on smaller data.	Data-efficient Transformers
Swin Transformer	Hierarchical ViT	Computes self-attention within shifted local windows for multi-scale dense prediction.	Object detection & segmentation
DINO / DINOv2 / DINOv3	Self-Supervised ViT	Self-distillation without labels producing rich generic visual features.	Foundation visual embeddings

4.2 Object Detection Architectures

Detector	Architecture Type	Mental Model & Mechanism	Primary Advantage
SSD	One-Stage CNN	Evaluates anchor boxes at multiple conv feature map resolutions in one pass.	Fast single-shot detection

Detector	Architecture Type	Mental Model & Mechanism	Primary Advantage
YOLO (v1-v10)	One-Stage CNN	Frames detection as a single regression grid problem for real-time inference.	Extreme speed & real-time video
Faster R-CNN	Two-Stage CNN	Uses Region Proposal Network (RPN) first, then crops and classifies candidate regions.	High spatial detection accuracy
RetinaNet	One-Stage CNN	Combines Feature Pyramid Network (FPN) with Focal Loss to handle severe background imbalance.	Accurate single-stage detection
FCOS	Anchor-Free CNN	Eliminates anchor boxes by predicting bounding box offsets directly at center points.	Anchor-free simplicity
DETR	Transformer Detector	Treats detection as a direct set prediction problem using Transformer encoder-decoder.	End-to-end NMS-free detection
Deformable DETR	Transformer Detector	Focuses self-attention on a small set of key sampling locations around reference points.	Fast multi-scale convergence
RT-DETR	Real-Time Transformer	First real-time end-to-end Transformer object detector balancing speed and accuracy.	Real-time Transformer detection
EfficientDet	Efficient Detector	Combines EfficientNet backbone with BiFPN multi-scale feature fusion.	Resource-efficient detection

4.3 Image Segmentation Architectures

Model	Segmentation Type	Mental Model & Mechanism	Best Used For
FCN	Semantic	Replaces fully-connected layers with deconvolutions for dense pixel classification.	First semantic segmentation
U-Net	Semantic	Symmetric U-shaped Encoder-Decoder with skip connections preserving spatial details.	Medical & small dataset masks
U-Net++	Semantic	Enhances U-Net with nested dense skip pathways to bridge semantic gap.	High-precision medical masks
DeepLabV3+	Semantic	Uses Atrous (dilated) Spatial Pyramid Pooling (ASPP) to capture multi-scale context.	Complex urban scene segmentation

Model	Segmentation Type	Mental Model & Mechanism	Best Used For
PSPNet	Semantic	Pyramid Scene Parsing network aggregating global context across spatial sub-regions.	Large context scene parsing
Mask R-CNN	Instance	Extends Faster R-CNN by adding a parallel branch for predicting pixel-level object masks.	Individual instance segmentation
YOLOACT	Real-Time Instance	Generates prototype masks and predicts linear coefficients for rapid real-time masks.	Fast real-time instance masks
SegFormer	Semantic Transformer	Lightweight MLP decoder combined with hierarchical Transformer encoder.	Efficient Transformer masks

Chapter 5: Unsupervised Learning & Vector Mathematics

5.1 Distance & Similarity Metrics

Metric	Formula / Concept	Mental Model	Best Application
Cosine Similarity	$\cos(\theta) = (A \cdot B) / (A B)$	Measures vector directional alignment	Text & visual vector search
Dot Product	$A \cdot B = \sum A_i \cdot B_i$	Measures magnitude and direction alignment	Dense embedding retrieval
Euclidean Distance	$d = \sqrt{(\sum (A_i - B_i)^2)}$	Straight-line physical distance between points	Standard spatial distance
Manhattan Distance	$d = \sum A_i - B_i $	Grid city-block distance along axes	High-dimensional sparse data
Minkowski Distance	$d = (\sum A_i - B_i ^p)^{1/p}$	Generalized distance family (p=1 Manhattan, p=2 Euclidean)	Parametric distance optimization
Mahalanobis Distance	$d = \sqrt{(x - \mu)^T \Sigma^{-1} (x - \mu)}$	Distance accounting for feature correlations	Outlier detection in multivariate data
Chebyshev Distance	$d = \max_i A_i - B_i $	Maximum absolute difference along any single axis	Chess king movement metric
Hamming Distance	$d = \sum I(A_i \neq B_i)$	Counts number of differing positions	Binary codes & string matching

5.2 Clustering Algorithms

Algorithm	Cluster Type	Mental Model & Mechanism	When to Use?
K-Means	Centroid-Based	Assigns points to nearest K centroids and updates centroids iteratively.	Known K, spherical clusters
Spherical K-Means	Directional Centroid	Cluster normalized vectors using cosine similarity.	Text & visual embeddings
DBSCAN	Density-Based	Grows clusters from dense neighborhoods; labels sparse points as noise.	Arbitrary shapes + noise filtering
HDBSCAN	Hierarchical Density	Extracts flat clusters from density hierarchy across varying densities.	Varying density, unknown K
OPTICS	Density Ordering	Orders points to analyze reachability plot across multiple density scales.	Complex density structures
Agglomerative	Hierarchical	Bottom-up merging of nearest clusters into a hierarchical tree (dendrogram).	Taxonomy & hierarchical groupings
Spectral Clustering	Graph-Based	Computes Laplacian matrix eigenvectors and clusters in low-dimensional space.	Complex non-convex cluster shapes
Mean Shift	Centroid Shift	Shifts points towards local modes of kernel density estimate.	Blob detection, unknown K

Algorithm	Cluster Type	Mental Model & Mechanism	When to Use?
GMM	Probabilistic	Fits mixture of K Multivariate Gaussians; provides soft probability membership.	Overlapping soft clusters

5.3 Dimensionality Reduction Techniques

Technique	Type	Mental Model & Mechanism	Primary Purpose
PCA	Linear Unsupervised	Rotates coordinate axes to directions of maximum variance.	Compression & preprocessing
LDA	Linear Supervised	Projects data to maximize between-class over within-class variance.	Supervised dimensionality reduction
t-SNE	Non-linear Manifold	Converts distances to probabilities and preserves local neighbor structures.	2D/3D visual cluster exploration
UMAP	Non-linear Manifold	Uses Riemannian geometry to preserve local + global manifold structure.	Fast 2D/3D visualization & ML input
Kernel PCA	Non-linear Kernel	Applies Kernel trick $K(x, z)$ before computing principal components.	Non-linear feature extraction
Truncated SVD	Linear Matrix	Performs SVD without centering data; works on sparse matrices.	Text TF-IDF & sparse matrices
Autoencoder	Neural Network	Encoder squeezes input into bottleneck code; decoder reconstructs.	Deep non-linear feature compression